

## Assignment No. 1

**Problem Statement:** Exploring data analysis (Various operations on dataset).

**Objective:** To perform Exploratory Data Analysis (EDA) and Preprocessing on a dataset to understand its structure, detect anomalies, and prepare it for machine learning models. The process includes handling missing data, analyzing correlations, applying encoding techniques, and visualizing data using charts and heatmaps.

### Prerequisite :

1. A Python environment set up with libraries like pandas, xml.etree.ElementTree, and requests (for web access).
2. Internet connection (for reading datasets from the web).
3. Text editor and basic knowledge of python and EDA

### Theory :

#### Exploratory Data Analysis (EDA) and Preprocessing

To build a well-performing machine learning model, it is essential to thoroughly explore and preprocess the dataset. The following steps ensure data quality, mitigate potential issues, and prepare the dataset for modeling.

##### 1. Understanding the Dataset

Before proceeding with any modifications, it is important to analyze the dataset's structure and characteristics. This helps identify potential inconsistencies and decide on necessary preprocessing techniques.

- **Dataset Dimensions**

1. The `.shape` function provides the total number of rows (samples) and columns (features).
2. If the dataset is too large, feature selection techniques might be required to prevent overfitting. Conversely, smaller datasets may need augmentation strategies.

- **Data Types of Columns**

1. Each column may contain numerical (integer/float) or categorical (string/object) values.
2. The `.info()` function provides a summary of the data types, helping determine whether encoding is necessary.

- **Missing Values**

1. Missing data can introduce biases, affecting model accuracy.
2. The `.isnull().sum()` function helps count missing values per column, indicating whether imputation or removal is needed.

- **Basic Statistical Measures**

1. Summary statistics such as mean, median, and standard deviation (via `.describe()`) provide an overview of data distribution.
2. If distributions are skewed, applying transformations like log-scaling may be beneficial.

## 2. **Handling Missing Data**

Missing values should be addressed to ensure data completeness and prevent biased model training. Common approaches include:

- **Removing Missing Data**

1. If a feature contains more than 50-60% missing values, it may be dropped due to insufficient information.
2. Rows with missing values can also be removed if their count is minimal and does not significantly impact the dataset.

- **Imputation Techniques**

1. **Numerical Data:**

- i. Mean imputation (for normally distributed data).
- ii. Median imputation (for skewed data).

2. **Categorical Data:**

- i. Mode imputation (replacing with the most frequent category).

## 3. **Correlation Analysis**

Analyzing relationships between numerical features helps identify redundant variables, reducing the risk of multicollinearity.

- **Methods to Analyze Correlation:**

- **Pearson's Correlation Coefficient:**

1. Measures the strength of linear relationships between numerical variables.
2. Values range from **-1 to +1**, where:
  - ii. **+1** indicates a strong positive correlation.
  - iii. **-1** indicates a strong negative correlation.
  - iv. **0** means no correlation.

- **Heatmap Visualization:**

1. A heatmap visually highlights highly correlated features, helping determine which ones to remove or merge.

#### 4. **Encoding Categorical Features**

Since machine learning models operate on numerical data, categorical variables must be converted appropriately.

- **Common Encoding Techniques:**

- **Label Encoding:**

- i. Assigns a unique integer to each category.
- ii. Suitable for ordinal variables (e.g., low < medium < high).

- **One-Hot Encoding (OHE):**

- i. Creates separate binary columns for each category.
- ii. Ideal for nominal variables (e.g., gender, city names).

#### 5. **Data Visualization**

Visualization helps in understanding patterns, distributions, and relationships within the data.

- **Commonly Used Plots:**

- **Histograms:** Show the frequency distribution of numerical variables.
- **Boxplots:** Help detect outliers in the dataset.
- **Scatter Plots:** Illustrate relationships between two numerical variables.

## 6. Feature Scaling and Normalization

Scaling numerical features ensures uniformity and improves model performance.

- **Scaling Techniques:**

- **Standardization (Z-score Normalization)**

- Converts data to a standard distribution with zero mean and unit variance.
- Useful for models like **linear regression, logistic regression, and PCA**.
- **Formula:**

$$X' = \frac{X - \mu}{\sigma}$$

- **Min-Max Scaling**

- Scales values between 0 and 1.
- Suitable for models such as **KNN and neural networks**.
- **Formula:**

$$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$$

- **Robust Scaling**

- Uses the median and interquartile range (IQR) to handle outliers.
- Recommended for datasets containing extreme values.
- **Formula:**

$$X' = \frac{X - \text{Median}}{\text{IQR}}$$

## Code & Output :

```
import pandas as pd
df= pd.read_csv("C:/Users/vishal/MachineLearning/Datasets/Titanic-Dataset.csv")
```

```
print(df.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

		Name	Sex	Age	SibSp	\
0		Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	1	
2		Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)		female	35.0	1	
4		Allen, Mr. William Henry	male	35.0	0	

	Parch		Ticket	Fare	Cabin	Embarked
0	0	A/5	21171	7.2500	NaN	S
1	0	PC	17599	71.2833	C85	C
2	0	STON/O2.	3101282	7.9250	NaN	S
3	0		113803	53.1000	C123	S
4	0		373450	8.0500	NaN	S

```
# shape of the data
df.shape
```

```
(891, 12)
```

```
df.tail(10)
```

```
#data information  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   PassengerId  891 non-null    int64  
1   Survived     891 non-null    int64  
2   Pclass       891 non-null    int64  
3   Name         891 non-null    object  
4   Sex          891 non-null    object  
5   Age          714 non-null    float64  
6   SibSp        891 non-null    int64  
7   Parch        891 non-null    int64  
8   Ticket       891 non-null    object  
9   Fare         891 non-null    float64  
10  Cabin        204 non-null    object  
11  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB
```

```
# describing the data  
df.describe()
```

```
Corr_Matrix = round(df.select_dtypes(include=[float, int]).corr(), 2)  
print(Corr_Matrix)
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
PassengerId	1.00	-0.01	-0.04	0.04	-0.06	-0.00	0.01
Survived	-0.01	1.00	-0.34	-0.08	-0.04	0.08	0.26
Pclass	-0.04	-0.34	1.00	-0.37	0.08	0.02	-0.55
Age	0.04	-0.08	-0.37	1.00	-0.31	-0.19	0.10
SibSp	-0.06	-0.04	0.08	-0.31	1.00	0.41	0.16
Parch	-0.00	0.08	0.02	-0.19	0.41	1.00	0.22
Fare	0.01	0.26	-0.55	0.10	0.16	0.22	1.00


```

import matplotlib.pyplot as plt
import seaborn as sns

axis_corr = sns.heatmap(
    Corr_Matrix,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(50, 500, n=500),
    square=True
)

plt.show()

```




```

import seaborn as sns
import matplotlib.pyplot as plt

# Distribution of Age
plt.figure(figsize=(7, 5))
sns.histplot(df['Age'], kde=True, bins=30)
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()

# Distribution of Fare
plt.figure(figsize=(7, 5))
sns.histplot(df['Fare'], kde=True, bins=30)
plt.title('Distribution of Fare')
plt.xlabel('Fare')
plt.ylabel('Frequency')
plt.show()


```





```
# Correlation Matrix (only numerical columns)
corr_matrix = df.select_dtypes(include=[float, int]).corr()




plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()
```



```
# Count plot for Sex
plt.figure(figsize=(7, 5))
sns.countplot(x='Sex', data=df)
plt.title('Count of Sex')
plt.xlabel('Sex')
plt.ylabel('Count')
plt.show()

# Count plot for Embarked
plt.figure(figsize=(7, 5))
sns.countplot(x='Embarked', data=df)
plt.title('Count of Embarked')
plt.xlabel('Embarked')
plt.ylabel('Count')
plt.show()


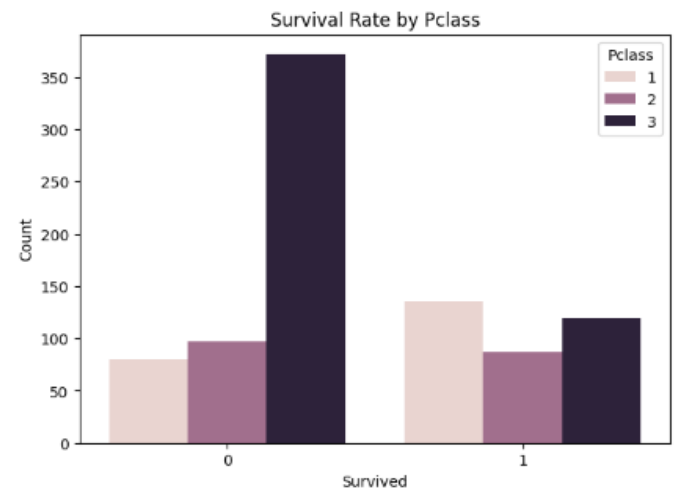

# Count plot for Pclass
plt.figure(figsize=(7, 5))
sns.countplot(x='Pclass', data=df)
plt.title('Count of Pclass')
plt.xlabel('Pclass')
plt.ylabel('Count')
plt.show()
```



```
# Survival Rate by Sex
plt.figure(figsize=(7, 5))
sns.countplot(x='Survived', hue='Sex', data=df)
plt.title('Survival Rate by Sex')
plt.xlabel('Survived')
plt.ylabel('Count')
plt.show()

# Survival Rate by Pclass
plt.figure(figsize=(7, 5))
sns.countplot(x='Survived', hue='Pclass', data=df)
plt.title('Survival Rate by Pclass')
plt.xlabel('Survived')
plt.ylabel('Count')
plt.show()

# Survival Rate by Embarked
plt.figure(figsize=(7, 5))
sns.countplot(x='Survived', hue='Embarked', data=df)
plt.title('Survival Rate by Embarked')
plt.xlabel('Survived')
plt.ylabel('Count')
plt.show()
```



```
# sum of missing values:  
df.isnull().sum()
```

```
PassengerId    0  
Survived        0  
Pclass         0  
Name           0  
Sex            0  
Age           177  
SibSp          0  
Parch          0  
Ticket         0  
Fare           0  
Cabin         687  
Embarked       2  
dtype: int64
```

```
# Calculate the percentage of missing values for each column  
missing_percentage = df.isnull().mean() * 100  
print(missing_percentage)
```

```
PassengerId    0.000000  
Survived        0.000000  
Pclass         0.000000  
Name           0.000000  
Sex            0.000000  
Age           19.865320  
SibSp          0.000000  
Parch          0.000000  
Ticket         0.000000  
Fare           0.000000  
Cabin         77.104377  
Embarked       0.224467  
dtype: float64
```

```
#checking duplicate values  
df.nunique()
```

```
PassengerId    891  
Survived        2  
Pclass         3  
Name           891  
Sex            2  
Age            88  
SibSp          7  
Parch          7  
Ticket        681  
Fare           248  
Cabin         147  
Embarked       3  
dtype: int64
```

```
# Fill missing values in 'Age' with the median of the column  
df['Age'] = df['Age'].fillna(df['Age'].median())
```

```
# Drop 'Cabin' as it has too many missing values and we don't have enough data to fill them  
df.drop(columns=['Cabin'], inplace=True, errors='ignore')
```

```
# Fill missing values in 'Embarked' with the mode of the column  
df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])
```

```

# Convert categorical columns 'Sex', 'Embarked' into numerical format using get_dummies
df = pd.get_dummies(df, columns=['Sex', 'Embarked'], drop_first=True)

# Drop non-feature columns
X = df.drop(columns=['Survived', 'Name', 'Ticket', 'PassengerId'])
y = df['Survived']

# Split the dataset into 70% training and 30% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Training and testing
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

model = LogisticRegression(max_iter=1000)

# Train the model on the training data
model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print(f"Confusion Matrix:\n{conf_matrix}")
print(f"Classification Report:\n{class_report}")

```


```

#Confusion Matrix Heatmap
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

conf_matrix = np.array([[136, 21], [30, 81]])
class_names = ['0', '1']

plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt='g', cmap='Blues', xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



```
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler

# Standardization
scaler_standard = StandardScaler()
df[['Age', 'Fare']] = scaler_standard.fit_transform(df[['Age', 'Fare']])

# Min-Max Scaling
scaler_minmax = MinMaxScaler()
df[['Age', 'Fare']] = scaler_minmax.fit_transform(df[['Age', 'Fare']])

# Robust Scaling
scaler_robust = RobustScaler()
df[['Age', 'Fare']] = scaler_robust.fit_transform(df[['Age', 'Fare']])
```

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# 1. Check Descriptive Statistics
print("Descriptive Statistics after Scaling:\n", df[['Age', 'Fare']].describe())

# 2. Check Mean and Standard Deviation
print("\nMean Values:\n", df[['Age', 'Fare']].mean())
print("\nStandard Deviation:\n", df[['Age', 'Fare']].std())

# 3. Check Minimum and Maximum Values
print("\nMinimum Values:\n", df[['Age', 'Fare']].min())
print("\nMaximum Values:\n", df[['Age', 'Fare']].max())

# 4. Check Data Distribution Using Histograms
df[['Age', 'Fare']].hist(figsize=(8, 4), bins=20)
plt.suptitle("Histograms of Scaled Features")
plt.show()

# 5. Check Outliers Using Boxplots
plt.figure(figsize=(8, 4))
sns.boxplot(data=df[['Age', 'Fare']])
plt.xticks(rotation=90)
plt.title("Boxplot of Scaled Features")
plt.show()

```

Descriptive Statistics after Scaling:

	Age	Fare
count	891.000000	891.000000
mean	0.104737	0.768745
std	1.001515	2.152200
min	-2.121538	-0.626005
25%	-0.461538	-0.283409
50%	0.000000	0.000000
75%	0.538462	0.716591
max	4.000000	21.562738

Mean Values:

Age 0.104737  
 Fare 0.768745  
 dtype: float64

Standard Deviation:

Age 1.001515  
 Fare 2.152200  
 dtype: float64

Minimum Values:

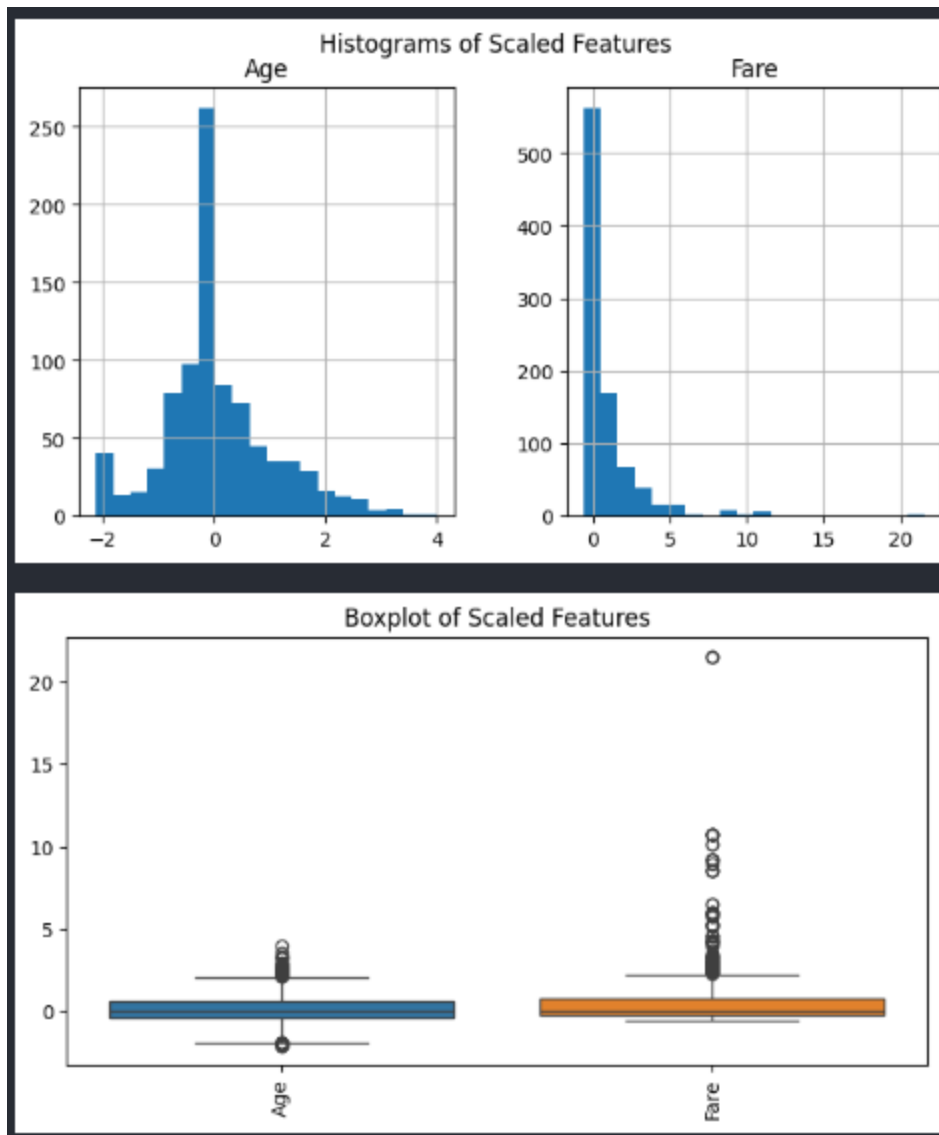
Age -2.121538  
 Fare -0.626005  
 dtype: float64

...

Maximum Values:

Age 4.000000  
 Fare 21.562738  
 dtype: float64

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings...](#)



**Github :** <https://github.com/Vishalgodalkar>

### **Conclusion:**

This EDA process involved inspecting the dataset, addressing missing values, encoding categorical variables, analyzing correlations, and applying feature scaling. Missing data was managed effectively, and duplicate features were identified. However, improper overwriting led to data corruption during scaling. To ensure accurate results, careful application of transformations is required. Overall, the preprocessing steps significantly improved data quality and prepared the dataset for machine learning applications.