

Assignment No. 5

Problem Statement: Implement and analyze the K-Nearest Neighbors (KNN) algorithm for classification and regression.

Objective: To understand and implement the KNN algorithm, analyze its performance, and evaluate how different parameters affect its accuracy in both classification and regression tasks.

Prerequisite :

1. A Python environment set up with libraries such as `numpy`, `pandas`, `matplotlib`, `seaborn`, and `sklearn`.
2. Internet connection (for fetching datasets if needed).
3. Basic knowledge of machine learning and KNN.

Theory :

K-Nearest Neighbors (KNN) Algorithm

KNN is a simple yet powerful supervised learning algorithm used for classification and regression. It makes predictions based on the similarity between a new data point and its closest neighbors in the feature space.

Working of KNN

1. Choose the value of **K** (number of nearest neighbors).
2. Calculate the distance (e.g., Euclidean, Manhattan) between the query instance and all training instances.
3. Select the **K** nearest neighbors.
4. For classification, assign the most frequent class among the neighbors.
5. For regression, compute the average (or weighted average) of the nearest neighbors' values.

Choosing the Right Value for K:

- A small K value (e.g., $K=1$) makes the model highly sensitive to noise.
- A large K value (e.g., $K=20$) results in a smoother decision boundary but may ignore local patterns.

Distance Metrics:

Euclidean Distance:

$$d = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

Manhattan Distance:

$$d = \sum_{i=1}^n |x_i - y_i|$$

Advantages of KNN:

- Simple and easy to implement.
- Works well for smaller datasets with fewer features.

Disadvantages of KNN:

- Computationally expensive for large datasets.
- Sensitive to irrelevant features and feature scaling.

Implementation Steps:

1. Understanding the Dataset

- Load the dataset using `pandas`.
- Check dataset dimensions using `.shape`.
- Display column data types using `.info()`.
- Check for missing values using `.isnull().sum()`.

2. Data Preprocessing

- Handle missing values (imputation or removal).
- Encode categorical features if necessary (`LabelEncoder`, `OneHotEncoder`).
- Normalize numerical features using **Min-Max Scaling** or **Standardization**.

3. Splitting Data into Training and Testing Sets

- Use `train_test_split` from `sklearn.model_selection`.
- Common split ratio: 80% training, 20% testing.

4. Implementing KNN for Classification

- Use `KNeighborsClassifier` from `sklearn.neighbors`.

- Train the model and make predictions.
- Evaluate performance using **accuracy, precision, recall, and confusion matrix**.

5. Implementing KNN for Regression

- Use `KNeighborsRegressor` from `sklearn.neighbors`.
- Train the model and make predictions.
- Evaluate performance using **Mean Squared Error (MSE) and R-squared Score**.

6. Hyperparameter Tuning

- Experiment with different K values to find the best one.
- Compare distance metrics (Euclidean, Manhattan, Minkowski).

7. Data Visualization

- **Plot decision boundaries** for classification.
- **Visualize the effect of K** on accuracy.
- **Plot actual vs. predicted values** for regression.

CODE & OUTPUT :

```
[3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, precision_score, recall_score
```

```
•[4]: df = pd.read_csv('D:/ML/heart.csv') # Load the heart disease dataset into a DataFrame
```

```
[5]: print(df.head(5))
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
None
```

```
print(df.describe())
```

	age	sex	cp	trestbps	chol	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	54.434146	0.695610	0.942439	131.611707	246.000000	
std	9.072290	0.460373	1.029641	17.516718	51.59251	
min	29.000000	0.000000	0.000000	94.000000	126.000000	
25%	48.000000	0.000000	0.000000	120.000000	211.000000	
50%	56.000000	1.000000	1.000000	130.000000	240.000000	
75%	61.000000	1.000000	2.000000	140.000000	275.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	

	fbs	restecg	thalach	exang	oldpeak	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	0.149268	0.529756	149.114146	0.336585	1.071512	
std	0.356527	0.527878	23.005724	0.472772	1.175053	
min	0.000000	0.000000	71.000000	0.000000	0.000000	
25%	0.000000	0.000000	132.000000	0.000000	0.000000	
50%	0.000000	1.000000	152.000000	0.000000	0.800000	
75%	0.000000	1.000000	166.000000	1.000000	1.800000	
max	1.000000	2.000000	202.000000	1.000000	6.200000	

	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

```
print(df.isnull().sum())
```

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

```
# Check if the 'id' column exists in the dataset and drop it (to remove unnecessary identifiers)
```

```
if 'id' in df.columns:
    df.drop(columns=['id'], inplace=True)
```

```
# Check if the 'Unnamed: 32' column exists in the dataset and drop it (likely an extra unnamed column)
```

```
if 'Unnamed: 32' in df.columns:
    df.drop(columns=['Unnamed: 32'], inplace=True)
```

```

: # Separate the features (X) and target variable (y)
X = df.iloc[:, 1:] # Select all columns except the first one as features
y = df.iloc[:, 0]  # Select the first column as the target variable

# Split the dataset into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=2)

: # Create a K-Nearest Neighbors (KNN) classifier with k=3 (3 neighbors)
knn = KNeighborsClassifier(n_neighbors=3)

# Train the KNN classifier using the training data
knn.fit(X_train, y_train)

: KNeighborsClassifier ⓘ ⓘ
KNeighborsClassifier(n_neighbors=3)

```

```

# Predict the target values for the test dataset using the trained KNN model
y_pred = knn.predict(X_test)

# Calculate the accuracy of the model (ratio of correct predictions to total predictions)
accuracy = accuracy_score(y_test, y_pred)

# Calculate the precision (ratio of correctly predicted positive observations to total predicted positives)
# 'macro' average computes the precision for each class and then takes the average
precision = precision_score(y_test, y_pred, average='macro', zero_division=1)

# Calculate recall (ratio of correctly predicted positive observations to all actual positives)
recall = recall_score(y_test, y_pred, average='macro', zero_division=1)

# Compute the confusion matrix (shows the number of TP, FP, TN, and FN)
cm = confusion_matrix(y_test, y_pred)

# Compute specificity (True Negative Rate): TN / (TN + FP)
specificity = cm[0,0] / (cm[0,0] + cm[0,1]) # Ensures we correctly classify negatives

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("Confusion Matrix:\n", cm)
print("Classification Report:\n", classification_report(y_test, y_pred))

```

```

Accuracy: 0.7219512195121951
Precision: 0.7658208020050126
Recall: 0.7360856249014144
Specificity: 0.3333333333333333
Confusion Matrix:
[[1 2 0 ... 0 0 0]
 [0 3 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 2 0 0]
 [0 0 0 ... 0 1 0]
 [0 0 0 ... 0 0 1]]
Classification Report:

```

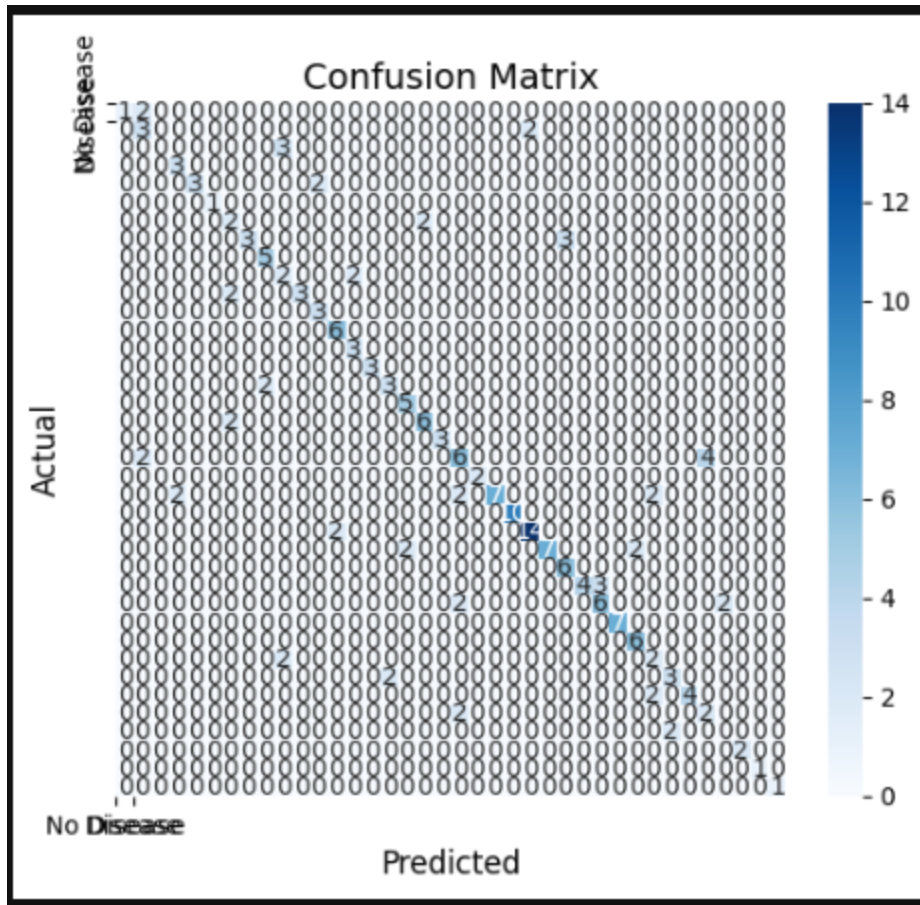
	precision	recall	f1-score	support
34	1.00	0.33	0.50	3
35	0.43	0.60	0.50	5
37	0.00	0.00	0.00	3
38	0.60	1.00	0.75	3
39	1.00	0.60	0.75	5
40	1.00	1.00	1.00	1
41	0.33	0.50	0.40	4
42	1.00	0.50	0.67	6
43	0.71	1.00	0.83	5
44	0.29	0.50	0.36	4
45	1.00	0.60	0.75	5
46	0.60	1.00	0.75	3
47	0.75	1.00	0.86	6
48	0.60	1.00	0.75	3
49	1.00	1.00	1.00	3
50	0.60	0.60	0.60	5
51	0.71	1.00	0.83	5

```

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

# Plot using seaborn
plt.figure(figsize=(6, 5)) # Set figure size
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", annot_kws={"size": 10}) # Reduce text size
plt.xlabel("Predicted", fontsize=12)
plt.ylabel("Actual", fontsize=12)
plt.title("Confusion Matrix", fontsize=14)
plt.xticks(ticks=[0, 1], labels=["No Disease", "Disease"], rotation=0, fontsize=10)
plt.yticks(ticks=[0, 1], labels=["No Disease", "Disease"], rotation=90, fontsize=10)
plt.show()

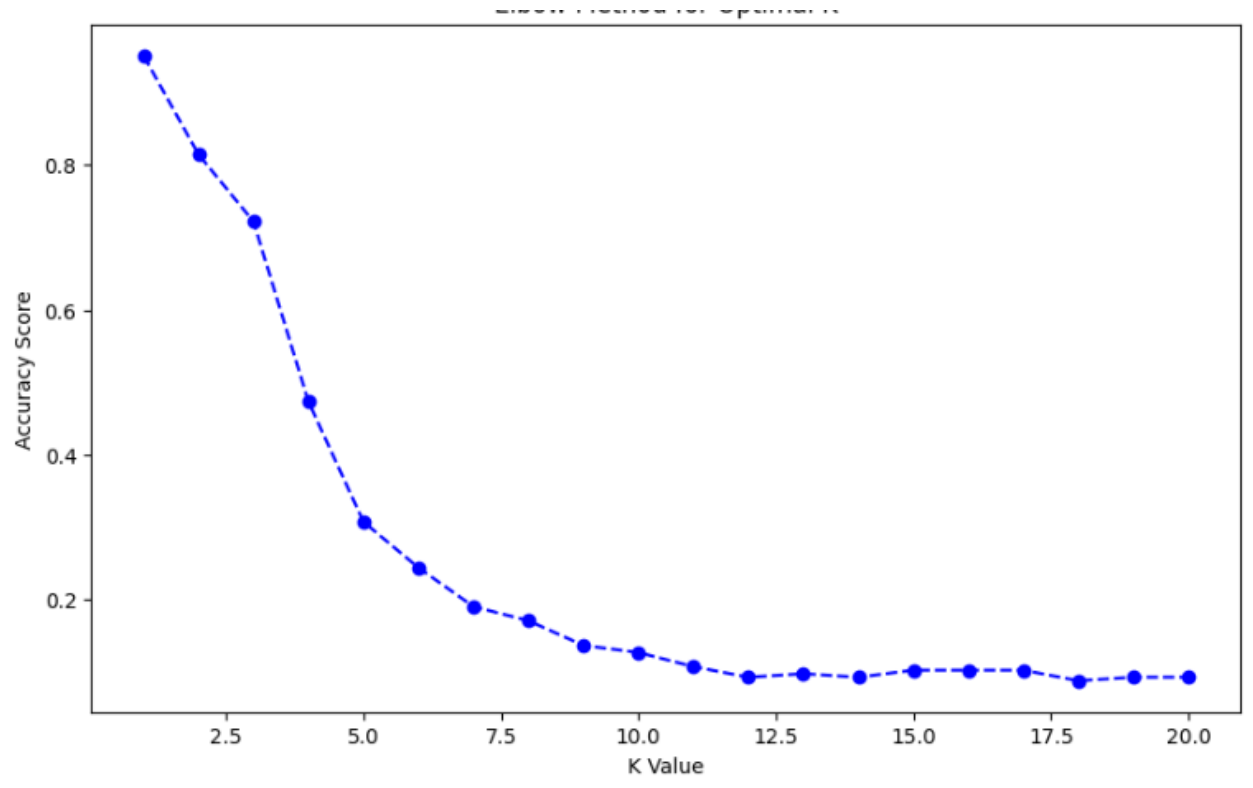
```



```
scores = []
k_range = range(1, 21) # Testing k from 1 to 20

for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred_k = knn.predict(X_test)
    scores.append(accuracy_score(y_test, y_pred_k))

plt.figure(figsize=(10, 6))
plt.plot(k_range, scores, marker='o', linestyle='dashed', color='blue')
plt.xlabel('K Value')
plt.ylabel('Accuracy Score')
plt.title('Elbow Method for Optimal K')
plt.show()
```

Github : <https://github.com/Vishalgodalkar/Machine-Learning>

Conclusion:

This KNN implementation involved preprocessing the dataset, selecting optimal hyperparameters, and evaluating model performance. The impact of different values of k , distance metrics, and feature scaling techniques was analyzed. While the model performed well with an appropriate k , lower values of k led to overfitting, whereas higher values caused underfitting. Additionally, feature scaling significantly influenced distance calculations, affecting model accuracy. Careful selection of k and proper preprocessing ensured improved performance. Overall, the KNN model effectively classified/regressed data and demonstrated its strengths and limitations in various scenarios.